

SQL Application Development

School of Computer Science
University of Waterloo

CS 348
Introduction to Database Management
Fall 2007

SQL APIs

- Interactive SQL command interpreters (e.g., DB2's command line processor) are simply domain-independent client programs that interact with an SQL database server
- In general, it is necessary to write other client programs for specific applications
- SQL has “bindings” for various programming languages that describe how applications written in those languages can be made to interact with a database server

Note

The main problem is the “impedance mismatch” between set-oriented SQL and the application programming language. How should data be passed back and forth between the two?

Outline

① Embedded SQL

- Static Embedded SQL

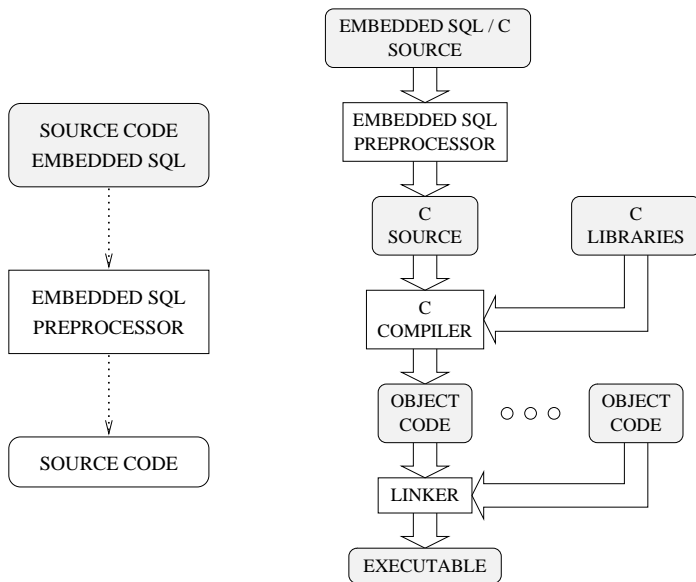
- Dynamic Embedded SQL

- SQLJ

② Call Level Interfaces

③ Stored Procedures

Development Process for Embedded SQL Applications



A Simple Example

```
#include <stdio.h>
EXEC SQL INCLUDE SQLCA;
main() {
    EXEC SQL WHENEVER SQLERROR GOTO error;
    EXEC SQL CONNECT TO sample;
    EXEC SQL UPDATE Employee
        SET salary = 1.1*salary
        WHERE empno = '000370';
    EXEC SQL COMMIT WORK;
    EXEC SQL CONNECT RESET;
    return(0);
error:
    printf("update failed, sqlcode = %ld\n",SQLCODE );
    return(-1);
}
```

Static Embedded SQL

- SQL DML and DDL can be embedded in a C program by prefixing with “EXEC SQL” and suffixing with “;”.
- host variables are used to send and receive values from the database system
 - values can be sent by using host variables in place of constants.
 - values can be received by using host variables in an INTO clause.

Note

The SELECT statement is (potentially) different in embedded SQL.

Declaring Host Variables

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
char deptno[4];
```

```
char deptname[30];
```

```
char mgrno[7];
```

```
char admrdept[4];
```

```
char location[17];
```

```
EXEC SQL END DECLARE SECTION;
```

```
/* program assigns values to variables */
```

```
EXEC SQL INSERT INTO
```

```
    Department(deptno,deptname,mgrno,admrcode,location)
```

```
VALUES
```

```
    (:deptno,:deptname,:mgrno,:admrcode,:location);
```

Domain and Type Correspondence

Domain	C Type
INTEGER	long int v;
SMALLINT	short int v;
REAL	float v;
DOUBLE	double v;
CHAR(<i>n</i>)	char v[n+1];
VARCHAR(<i>n</i>)	char v[n+1]; or struct tag { short int len; char v[n]; }
DATE	char v[11];

Note

Each SQL domain (type) corresponds to a type in the host language. See, e.g., the DB2 Application Development Guide for complete list.

Queries Using INTO

- Print the last name of a specified employee.

```
int PrintEmployeeName( char employeenum[] ) {  
EXEC SQL BEGIN DECLARE SECTION;  
    char empno[7];  
    char lastname[16];  
EXEC SQL END DECLARE SECTION;  
    strcpy(empno,employeenum);  
EXEC SQL  
        SELECT lastname INTO :lastname  
        FROM employee  
        WHERE empno = :empno;  
    if( SQLCODE < 0 ) { return( -1 ); } /* error */  
    else if(SQLCODE==100){printf("no such employee\n");}  
    else { printf("%s\n",lastname); }  
    return( 0 );  
}
```

Indicator Variables


- What if a returned value is NULL?
 - NULLs are handled using special flags called *indicator variables*.
 - Any host variable that might receive a NULL should have a corresponding indicator variable.
 - In C/C++, indicator variables are short ints

Indicator Variables: An Example

```
int PrintEmployeePhone( char employeenum[] ) {  
    EXEC SQL BEGIN DECLARE SECTION;  
        char empno[7];  
        char phonenum[5];  
        short int phoneind;  
    EXEC SQL END DECLARE SECTION;  
        strcpy(empno,employeenum);  
    EXEC SQL  
        SELECT phoneno INTO :phonenum :phoneind  
        FROM employee WHERE empno = :empno;  
    if( SQLCODE < 0 ) { return( -1 ); } /* error */  
    else if(SQLCODE==100){printf("no such employee\n");}  
    else if (phoneind<0){printf("phone unknown\n");}  
    else { printf("%s\n",phonenum); }  
    return( 0 );  
}
```

Cursors

- If a query may return more than one row, then a *cursor* must be used to retrieve values from the result.
- A cursor is like a pointer that refers to some row of the result. At any time, a cursor may be in one of three places:
 - before first tuple
 - on a tuple
 - after last tuple



Using Cursors

- 1 Declare the cursor
 - Declaring a cursor associates a cursor identifier with a query.
- 2 Open the cursor
 - Opening a cursor (conceptually) causes the query to be evaluated, generating a result.
- 3 Fetch one or more tuples using the cursor
 - Each call to the `FETCH` command returns values from one tuple of the generated result.
- 4 Close the cursor

The FETCH Command Syntax

```
fetch [<location>] <cursor-name>  
      [ INTO <host-var1>, <host-var2> ... ]
```

- Possible locations:
 - NEXT (this is the default)
 - PRIOR
 - FIRST
 - LAST
 - ABSOLUTE *n*
 - RELATIVE *n*

Using Cursors: An Example

```
int PrintEmpNames() {
    int rval; /* -1 for error, 0 for success */
    EXEC SQL BEGIN DECLARE SECTION;
    char fullname[30];
    EXEC SQL END DECLARE SECTION;
    EXEC SQL DECLARE C1 CURSOR FOR
        SELECT firstnme || ' ' || lastname FROM Employee;
    EXEC SQL OPEN C1;
    for( ;; ) {
        EXEC SQL FETCH NEXT C1 INTO :fullname;
        if (SQLCODE == 100) { rval = 0; break; }
        else if (SQLCODE < 0) { rval = -1; break; }
        printf("%s\n", fullname );
    }
    EXEC SQL CLOSE C1;
    return(rval); }
```

Dynamic Embedded SQL

- Must be used when tables, columns or predicates are not known at the time the application is written.
- Basic idea:
 - ① prepare the statement for execution: `PREPARE`
 - in static embedded SQL programs, statement preparation is handled at compile time by the preprocessor
 - ② execute the prepared statement: `EXECUTE`
- Once prepared, a statement may be executed multiple times, if desired

Dynamic Embedded SQL: A Simple Example

```
EXEC SQL BEGIN DECLARE SECTION;  
char s[100] =  
    "INSERT INTO department VALUES ('000456','Legal',...)"  
EXEC SQL END DECLARE SECTION;  
EXEC SQL EXECUTE IMMEDIATE :s;
```

or, to factor cost of “preparing”

```
EXEC SQL BEGIN DECLARE SECTION;  
char s[100] =  
    "INSERT INTO department VALUES ('000456','Legal',...)"  
EXEC SQL END DECLARE SECTION;  
EXEC SQL PREPARE S1 FROM :s;  
EXEC SQL EXECUTE S1;  
EXEC SQL EXECUTE S1;
```

Dynamic Embedded SQL: Using Host Variables for Input

```
EXEC SQL BEGIN DECLARE SECTION;
char s[100] = "INSERT INTO employee VALUES (?, ?, ...
char empno[7];
char firstname[13];
...
EXEC SQL END DECLARE SECTION;

EXEC SQL PREPARE S1 FROM :s;
strcpy(empno, "000111");
strcpy(firstname, "Ken");
...
EXEC SQL EXECUTE S1 USING :empno, :firstname, ... ;
```

Placeholders

- In the query string

```
INSERT INTO employee VALUES (?, ?, ... );
```

the ? are called *placeholders*

- placeholders can appear where literals can appear - not in place of relation names, column names, etc.
- host variable values replace the placeholders when the prepared statement is executed
- the USING clause is used to specify which host variables should replace the placeholders:

```
EXEC SQL EXECUTE S1 USING :empno, :firstname, ... ;
```
- USING can only be used with previously-prepared statements, not with EXECUTE IMMEDIATE

Dynamic Single-Row Queries

```
EXEC SQL BEGIN DECLARE SECTION;
char s[100] =
    "select lastname,salary from employee where empno = ?";
char empno[7];
char lastname[16];
double salary;
short int salaryind;
EXEC SQL END DECLARE SECTION;
EXEC SQL PREPARE S1 FROM :s;
EXEC SQL EXECUTE S1
    INTO :lastname, :salary:salaryind USING :empno
```

- INTO (with EXECUTE) in dynamic SQL is like INTO (with SELECT) in static
- Note: our DB2 version does not allow the use of INTO with EXECUTE. A dynamic cursor must be used to retrieve values.

Dynamic Cursors

```
EXEC SQL BEGIN DECLARE SECTION;
char s[100] =
    "select lastname,salary from employee where edlevel =
short int edlevel;
char lastname[16];
double salary;
short int salaryind;
EXEC SQL END DECLARE SECTION;
EXEC SQL PREPARE S1 FROM :s;
EXEC SQL DECLARE C1 CURSOR FOR S1;
edlevel = 18;
EXEC SQL OPEN C1 USING :edlevel;
while( ... ) {
    EXEC SQL FETCH FROM C1
        INTO :lastname, :salary:salaryind;
}
```

Descriptors and the SQLDA

- if the numbers and types of input and output values are not known in advance, SQL *descriptors* can be used determine them at run-time
- an SQLDA (descriptor area) is used to hold a description of the structure (number of attributes and their types) of a query result.
- the DESCRIBE command can be used to populate a descriptor area, that is, to find out the structure of a query result

- SQLJ allows embedding of SQL into Java
- Not part of SQL standard, but supported by most DBMSs
- Like Embedded SQL, utilizes preprocessing step
 - static type checking against database schema
 - DBMS can optimize static queries at compile time
- Unlike Embedded SQL, runtime connection established via JDBC connection
 - forces compliance to SQL standard syntax

JDBC, ODBC and CLI

- CLI (Call-Level Interface) is a vendor-neutral ISO standard programming interface for SQL database systems. It is similar to ODBC.
- ODBC (Open Database Connectivity), popularized by Microsoft, is a programming interface for SQL database systems.
- JDBC (Java Database Connectivity) is a collection of Java classes that provide an ODBC/CLI-like programming interface.
- An embedded SQL program used to access one DBMS must be recompiled before it can be used to access a different DBMS.
- A CLI/ODBC/JDBC program need not be recompiled - a single application may even access multiple DBMS at the same time.

A CLI Example

```
SQLHANDLE henv; /* an environment handle*/
SQLHANDLE hdbc; /* a connection handle */
SQLHANDLE hstmt; /* a statement handle */
SQLCHAR numteamsquery[] = "select count(*) from teams"
SQLAllocHandle(SQL_HANDLE_ENV, SQL_NULL_HANDLE, &henv)
DBconnect(henv,&hdbc,server,uid,pwd);
SQLAllocHandle( SQL_HANDLE_STMT, hdbc, &hstmt );
SQLExecDirect(hstmt,numteamsquery,SQL_NTS ); /* execut
SQLFetch(hstmt); /* get one row of the result */
SQLGetData(hstmt,1,SQL_C_LONG,&numteams,
           sizeof(numteams),&bytesremaining);
SQLFreeStmt(hstmt,SQL_CLOSE); /* close the statement
```

Note

CLI/ODBC interface is similar to dynamic embedded SQL, but syntax is entirely valid host language.

Stored Procedures

Idea

A stored procedure executes application logic directly inside the DBMS process.

- Possible implementations
 - invoke externally-compiled application
 - SQL/PSM (or vendor-specific language)
- Possible advantages of stored procedures:
 - 1 minimize data transfer costs
 - 2 centralize application code
 - 3 logical independence

A Stored Procedure Example: Atomic-Valued Function

```
CREATE FUNCTION deptTotalSalaries(dept CHAR(3))  
    RETURNS DECIMAL(9,2)  
LANGUAGE SQL  
RETURN  
    SELECT sum(salary)  
    FROM employee  
    WHERE workdept = dept
```

A Stored Procedure Example: Atomic-Valued Function

```
db2 => SELECT deptno, sumSalaries(deptno) AS sal \
=> FROM department
```

DEPTNO	SAL
A00	128500.00
B01	41250.00
C01	90470.00
D01	-
D11	222100.00
D21	150920.00
E01	40175.00
E11	104990.00
E21	95310.00

9 record(s) selected.

A Stored Procedure Example: Table-Valued Function

```
CREATE FUNCTION deptSalariesF(dept CHAR(3))  
    RETURNS TABLE(salary DECIMAL(9,2))  
    LANGUAGE SQL  
RETURN  
    SELECT salary  
    FROM employee  
    WHERE workdept = dept
```

A Stored Procedure Example: Table-Valued Function

```
db2 => SELECT * FROM TABLE \  
=> (deptSalariesF(CAST('A00' AS CHAR(3)))) AS s
```

SALARY

52750.00

46500.00

29250.00

3 record(s) selected.

A Stored Procedure Example: Multiple Results

```
CREATE PROCEDURE deptSalariesP(IN dept CHAR(3))  
    RESULT SETS 2  
    LANGUAGE SQL  
BEGIN  
    DECLARE emp_curs CURSOR WITH RETURN FOR  
        SELECT salary  
        FROM employee  
        WHERE workdept = dept;  
  
    DECLARE dept_curs CURSOR WITH RETURN FOR  
        SELECT deptno, sumSalaries(deptno) as sumsal  
        FROM department;  
  
    OPEN emp_curs;  
    OPEN dept_curs;  
  
END
```

A Stored Procedure Example: Multiple Results

```
db2 => call deptSalariesP('A00')
```

```
SALARY
```

```
52750.00
```

```
46500.00
```

```
29250.00
```

```
DEPTNO SUMSAL
```

```
A00      128500.00
```

```
B01      41250.00
```

```
C01      90470.00
```

```
D01      NULL
```

```
D11      222100.00
```

```
D21      150920.00
```

```
E01      40175.00
```

```
E11      104990.00
```

```
E21      95310.00
```

```
"DEPTSALARIESP" RETURN_STATUS: "0"
```


A Stored Procedure Example: Branching

```
CREATE PROCEDURE UPDATE_SALARY_IF
    (IN employee_number CHAR(6), INOUT rating SMALLINT)
    LANGUAGE SQL
BEGIN
    DECLARE not_found CONDITION FOR SQLSTATE '02000';
    DECLARE EXIT HANDLER FOR not_found
        SET rating = -1;
    IF rating = 1 THEN
        UPDATE employee
        SET salary = salary * 1.10, bonus = 1000
        WHERE empno = employee_number;
    ELSEIF rating = 2 THEN
        UPDATE employee
        SET salary = salary * 1.05, bonus = 500
        WHERE empno = employee_number;
    ELSE
        UPDATE employee
        SET salary = salary * 1.03, bonus = 0
        WHERE empno = employee_number;
    END IF;
END
```